

ROBOT SKILL HALF-LIFE UNDER SHIFT: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study continual robot learning. The motivating bet is: Quantify how quickly skills decay under small physical environment shifts. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For continual robot learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b; 2012; 2019; 2021; 2025). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.532	0.468
observed_only_wam	hidden_mode_shift	0.168	0.833
augmentation_as_data_wam	nominal	0.613	0.387
augmentation_as_data_wam	hidden_mode_shift	0.240	0.760
scalar_uncertainty_wam	nominal	0.647	0.353
scalar_uncertainty_wam	hidden_mode_shift	0.378	0.623
proposed_branch_mechanism	nominal	0.762	0.238
proposed_branch_mechanism	hidden_mode_shift	0.545	0.455

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Robot Skill Half-Life Under Shift argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Lost in translation? skills policy and the shift to skill ecosystems, 2012. <https://openalex.org/W161294835>.
- Learning and transfer of complex motor skills in virtual reality: a perspective review, 2019. <https://doi.org/10.1186/s12984-019-0587-8>.
- Transformative top-down planning in a small african city: How residents in bagamoyo, tanzania connect with a city in motion, 2021. <https://doi.org/10.1177/2399654419864605>.
- Evaluating how psychological senses and physical motions are affected by avatar shapes in a non-immersive environment, 2025. <https://doi.org/10.1007/s10015-024-00979-y>.
- Planning under distribution shifts with causal pomdps, 2026a. <http://arxiv.org/abs/2602.23545v2>.
- Uncertainty-calibrated safety gating for vision-language- action manipulation under domain shift: Reliability gains and intervention-efficiency trade-offs, 2026b. <https://doi.org/10.3390/s26103140>.