

ROBOT SKILL HALF-LIFE UNDER SHIFT: ICLR MAIN-CONFERENCE GATE ARCHIVE

Anonymous authors

Paper under double-blind review

ABSTRACT

This document is not an ICLR main-conference submission. It is the v3 main-conference gate for a generated robotics draft about continual robot learning. The original research bet was: Quantify how quickly skills decay under small physical environment shifts. After a stricter hostile review, the honest terminal decision is **KILL/ARCHIVE**. Submission-hardening version: v3.

1 GATE DECISION

The v2 paper was workshop-only. Under the active ICLR-main standard, that is insufficient. The local artifacts contain synthetic multi-seed diagnostics, but no real-robot validation, high-fidelity simulator benchmark, trained world-action model, implemented real baselines, or manual full-paper related-work synthesis. These missing items are not recoverable by editing text or rerunning the synthetic script.

2 FATAL ICLR-MAIN BLOCKERS

- Evidence is synthetic and generated from a shared batch template.
- Baselines are diagnostic probability models, not implemented competing robot systems.
- Ablations do not correspond to a trained model checkpoint.
- No hardware, high-fidelity benchmark, qualitative rollouts, or failure videos are available.
- Related work is based on pool slices and hostile metadata, not a manual full-paper synthesis.
- Claims would be misleading if framed as ICLR-main empirical robotics evidence.

3 HOSTILE PRIOR-WORK PRESSURE

The closest local hostile records include:

- Planning under Distribution Shifts with Causal POMDPs (2026)
- Uncertainty-Calibrated Safety Gating for Vision-Language- Action Manipulation Under Domain Shift: Reliability Gains and Intervention-Efficiency Trade-Offs (2026)
- Lost in translation? Skills policy and the shift to skill ecosystems (2012)
- Learning and transfer of complex motor skills in virtual reality: a perspective review (2019)
- Transformative top-down planning in a small African city: How residents in Bagamoyo, Tanzania connect with a city in motion (2021)

These works make generic world models, uncertainty, model-based planning, and robot policy robustness crowded. Without real empirical evidence, the remaining novelty boundary is too weak for ICLR main.

4 WHAT REMAINS USEFUL

The archive is still useful as an idea seed and reproducibility scaffold. The synthetic code, stress-test CSVs, reviewer attacks, and novelty-boundary docs can inform a future project. The correct next step is not more paper polishing; it is new evidence: implemented models, real baselines, manual related work, and robot or high-fidelity simulator validation.

5 TERMINAL ACTION

Decision: **KILL/ARCHIVE for ICLR main**. Do not submit this paper to ICLR main in the current form. Revival requires a substantially new empirical project.